

Build PyTorch on HPC Cluster

Bosung Lee

bslee@nextfoam.co.kr

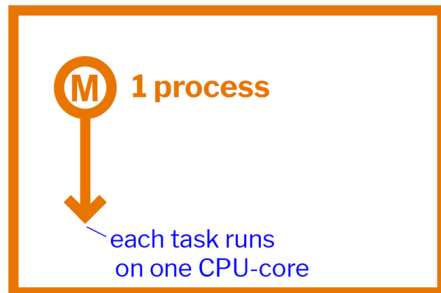
CPO, NEXTFOAM

Four Basic Types of Parallel Programming

Embarassingly Parallel

- each task can be run completely independently of other tasks
- no communication required between tasks

1 computer (node)

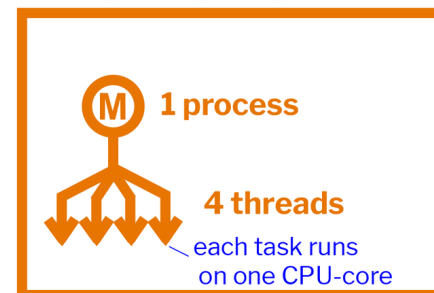


x 50

Shared-Memory Parallel

- tasks are run as threads on separate CPU-cores
- light level of communication between the cores working on each task

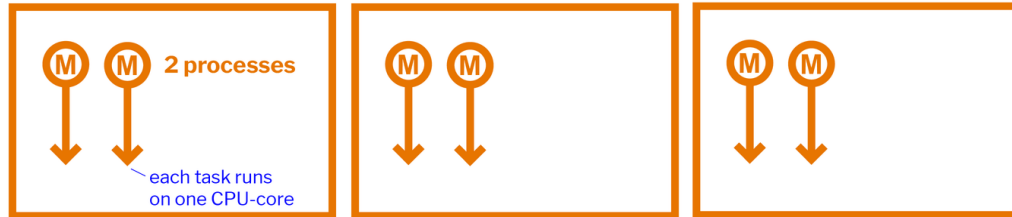
1 computer (node)



Distributed Memory Parallel

- multiple processes do not share the same space in memory
- high level of communication between different tasks

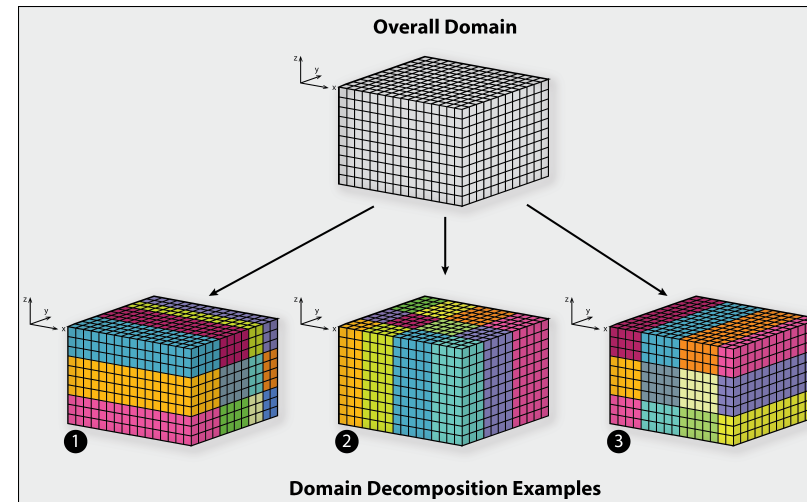
3 computers (nodes)



Accelerator Parallel

- uses different types of computer hardware - such as GPU and FPGA

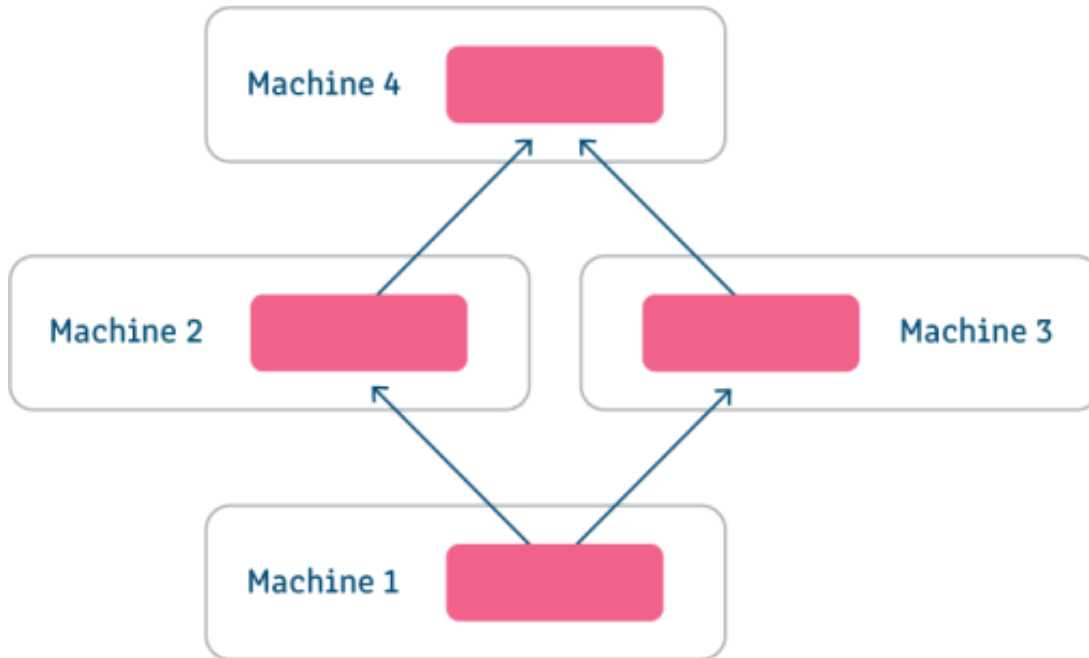
Domain Decomposition



Distributed Data Parallel vs Model Parallel

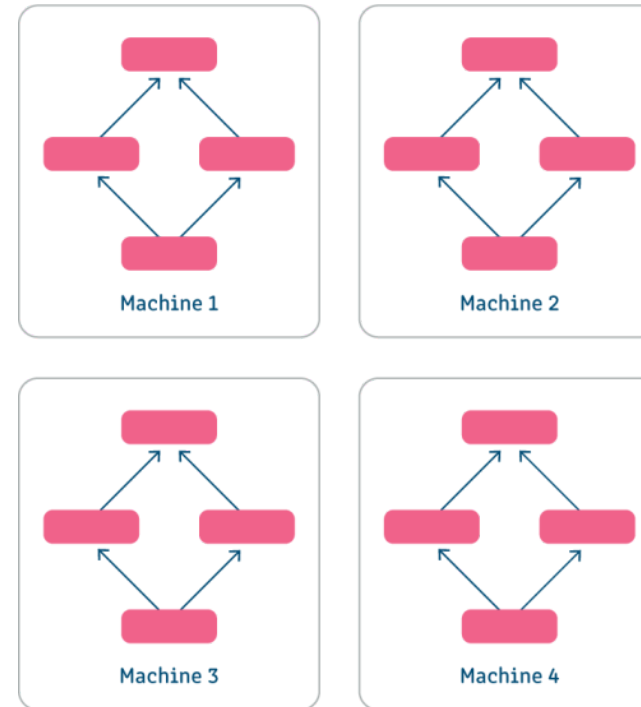
- the model itself is divided into parts that are trained simultaneously across different worker nodes

Model Parallelism



- Each worker node trains a copy of the model on a different batch of training data

Data Parallelism



Build HPC for PyTorch Distributed Data Parallel

1. HPC Cluster Setup
 - Configure communication backends (gloo, mpi, nccl)
2. Install [miniconda](#) for all compute nodes
3. [Install Pytorch with MPI Backend](#)
4. [Run distributed training using MPI backend](#)

Communication Backends

Backend	gloo		mpi		nccl	
Device	CPU	GPU	CPU	GPU	CPU	GPU
send	✓	✗	✓	?	✗	✓
recv	✓	✗	✓	?	✗	✓
broadcast	✓	✓	✓	?	✗	✓
all_reduce	✓	✓	✓	?	✗	✓
reduce	✓	✗	✓	?	✗	✓
all_gather	✓	✗	✓	?	✗	✓
gather	✓	✗	✓	?	✗	✓
scatter	✓	✗	✓	?	✗	✓
reduce_scatter	✗	✗	✗	✗	✗	✓
all_to_all	✗	✗	✓	?	✗	✓
barrier	✓	✗	✓	?	✗	✓

Which backend to use?

- Rule of thumb
 - Distributed GPU training : NCCL
 - Distributed CPU training. : Gloo
- GPU hosts with InfiniBand : NCCL
- GPU hosts with Ethernet : NCCL
- CPU hosts with InfiniBand : Gloo, MPI
- CPU hosts with Ethernet : Gloo, MPI

Install OpenMPI 4.1.6

Login to the headnode as root and install required packages to compile. Before installing the OpenMPI from source, remove default openmpi packages

```
root@hostname:~# apt remove openmpi*
root@headnode:~# apt-get -y update
root@headnode:~# apt-get -y install build-essential flex zlib1g-dev libgmp-dev libmpfr-dev
```

Download `openmpi 4.1.6` source and install to `/opt/openmpi-4.1.6` directory

```
root@headnode:~# wget https://download.open-mpi.org/release/open-mpi/v4.1/openmpi-4.1.6.tar.gz
root@headnode:~# tar xzf openmpi-4.1.6.tar.gz ; rm openmpi-4.1.6.tar.gz
root@headnode:~# cd openmpi-4.1.6
root@headnode:~# ./configure --prefix=/opt/openmpi-4.1.6
root@headnode:~# make -j 4 all
root@headnode:~# make install
root@headnode:~# echo 'export PATH=$PATH:/opt/openmpi-4.1.6/bin' >> /etc/bash.bashrc
```

For GPUs, `--with-cuda` or `--with-rocm` option is needed at `configure` step

Install miniconda for all compute nodes

Login to the headnode as root and install miniconda for all compute nodes

```
root@headnode:~# wget https://repo.anaconda.com/miniconda/Miniconda3-latest-Linux-x86_64.sh  
-O /opt/miniconda-installer.sh  
root@headnode:~# bash /opt/miniconda-installer.sh
```

Once the installation is finished, type `[yes]` to initialize Miniconda

```
Do you wish the installer to initialize Miniconda3  
by running conda init? [yes|no]  
[no] >>> yes
```

Set the `auto_activate_base` to `false` not to activate the base environment at startup

```
root@headnode:~# conda config --set auto_activate_base false
```


Add the conda initialization script at `/etc/bash.bashrc` for all users

```
# >>> conda initialize >>>
# !! Contents within this block are managed by 'conda init' !!
__conda_setup="$('/opt/miniconda3/bin/conda' 'shell.bash' 'hook' 2> /dev/null)"
if [ $? -eq 0 ]; then
    eval "$__conda_setup"
else
    if [ -f "/opt/miniconda3/etc/profile.d/conda.sh" ]; then
        . "/opt/miniconda3/etc/profile.d/conda.sh"
    else
        export PATH="/opt/miniconda3/bin:$PATH"
    fi
fi
unset __conda_setup
# <<< conda initialize <<<
```

Copy miniconda installation to all compute nodes

```
root@headnode:~# conda update -y --all
root@headnode:~# bpushdir /opt/miniconda3 /opt
root@headnode:~# bpush /etc/bash.bashrc /etc
```

Build Pytorch with MPI backend

To get MPI backend for `torch distributed` working you need to recompile PyTorch. To compile Pytorch with MPI, login as a user and create a virtual environment

```
bosung@node001:~$ conda create -n torch-mpi python=3.8  
bosung@node001:~$ conda activate torch-mpi
```

Download pytorch and install required packages, and compile

```
(torch-mpi) bosung@node001:~$ git clone https://github.com/pytorch/pytorch  
(torch-mpi) bosung@node001:~$ pip install pyyaml typing_extensions torchvision  
(torch-mpi) bosung@node001:~$ conda install -c conda-forge libstdcxx-ng=12  
(torch-mpi) bosung@node001:~$ cd pytorch  
(torch-mpi) bosung@node001:~/pytorch$ python3 setup.py develop --user
```

Run distributed training using MPI backend

Login to a head node and activate the virtual environment. Then clone the repo and download dataset to train a simple CNN on the MNIST data set.

```
bosung@node001:~$ conda activate torch-mpi
(torch-mpi) bosung@node001:~$ git clone https://github.com/PrincetonUniversity/install_pytorch.git
(torch-mpi) bosung@node001:~$ cd install_pytorch
(torch-mpi) bosung@node001:~/install_pytorch$ python download_mnist.py
```

To print the stdout, install `sys` and modify the source code

```
print('Train Epoch: {} [{}/{ } ({:.0f}%)]\tLoss: {:.6f}'.format(
    epoch, batch_idx * len(data), len(train_loader.dataset),
    100. * batch_idx / len(train_loader), loss.item()))
sys.stdout.flush()
```

Run the DDP training multi cores.

```
(torch-mpi) bosung@node001:~/install_pytorch$ mpirun -np 16 python mnist_classify.py --epochs=3
```